

Differentiable Programming

Atılım Güneş Baydin

National University of Ireland Maynooth

(Based on joint work with Barak Pearlmutter)

Microsoft Research Cambridge, February 1, 2016



Hamilton Institute



**Maynooth
University**
National University
of Ireland Maynooth

Deep learning layouts

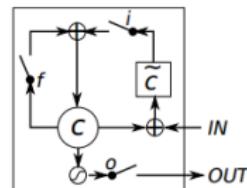
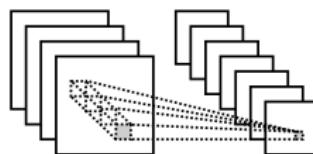
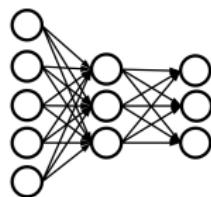
Neural network models are assembled from **building blocks** and trained with **backpropagation**

Deep learning layouts

Neural network models are assembled from **building blocks** and trained with **backpropagation**

Traditional:

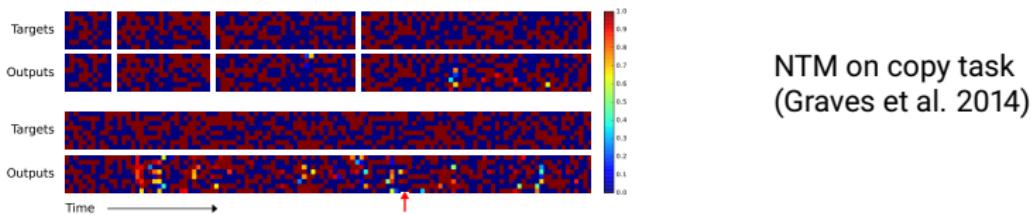
- Feedforward
- Convolutional
- Recurrent



Deep learning layouts

Newer additions:

Make **algorithmic** elements **continuous and differentiable**
→ enables use in deep learning



NTM on copy task
(Graves et al. 2014)

- Neural Turing Machine (Graves et al., 2014)
→ can infer algorithms: copy, sort, recall
- Stack-augmented RNN (Joulin & Mikolov, 2015)
- End-to-end memory network (Sukhbaatar et al., 2015)
- Stack, queue, deque (Grefenstette et al., 2015)
- Discrete interfaces (Zaremba & Sutskever, 2015)

Deep learning layouts

Stacking of many layers, trained through backpropagation

AlexNet, 8 layers (ILSVRC 2012)



VGG, 19 layers (ILSVRC 2014)



ResNet, 152 layers (deep residual learning) (ILSVRC 2015)



(He, Zhang, Ren, Sun. "Deep Residual Learning for Image Recognition." 2015. arXiv:1512.03385)

The bigger picture

One way of viewing deep learning systems is
“differentiable functional programming”

Two main characteristics:

- **Differentiability**

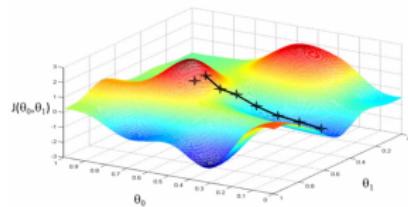
- optimization

- Chained function **composition**

- successive transformations

- successive levels of distributed representations (Bengio 2013)

- the chain rule of calculus propagates derivatives



$$g : A \rightarrow B$$

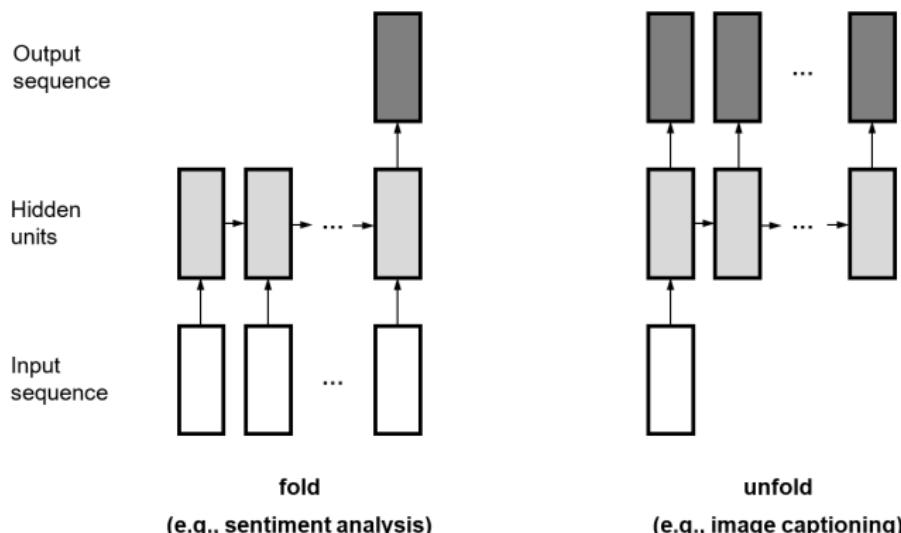
$$f : B \rightarrow C$$

$$f \circ g : A \rightarrow C$$

The bigger picture

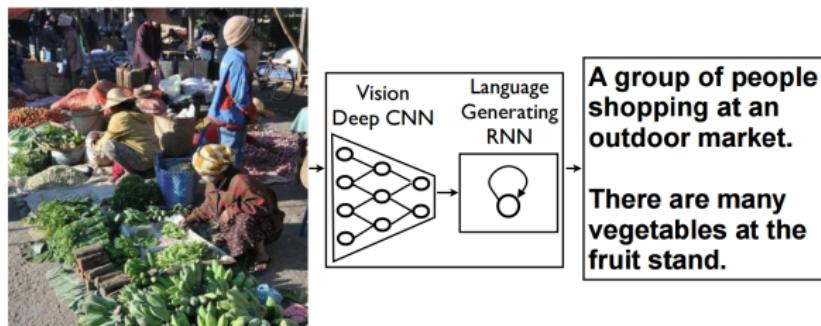
In a functional interpretation

- **Weight-tying** or multiple applications of the same neuron (e.g., ConvNets and RNNs) resemble **function abstraction**
- **Structural patterns** of composition resemble **higher-order functions** (e.g., map, fold, unfold, zip)



The bigger picture

Even when you have **complex compositions**, differentiability ensures that they can be trained end-to-end with backpropagation



(Vinyals, Toshev, Bengio, Erhan. "Show and tell: a neural image caption generator." 2014. arXiv:1411.4555)

The bigger picture

These insights clearly put into words in
Christopher Olah's blog post (September 3, 2015)

<http://colah.github.io/posts/2015-09-NN-Types-FP/>

"The field does not (yet) have a unifying insight or narrative"

and reiterated in David Dalrymple's essay (January 2016)

<http://edge.org/response-detail/26794>

"The most natural playground ... would be a new language that can run back-propagation directly on functional programs."

In this talk

Vision:

Functional languages with

- deeply embedded,
- general-purpose

differentiation capability, i.e., **differentiable programming**

In this talk

Vision:

Functional languages with

- deeply embedded,
- general-purpose

differentiation capability, i.e., **differentiable programming**

Automatic (algorithmic) differentiation (AD) in a functional framework is a manifestation of this vision.

In this talk

I will talk about:

- Mainstream frameworks
- What AD research can contribute
- My ongoing work

Mainstream Frameworks

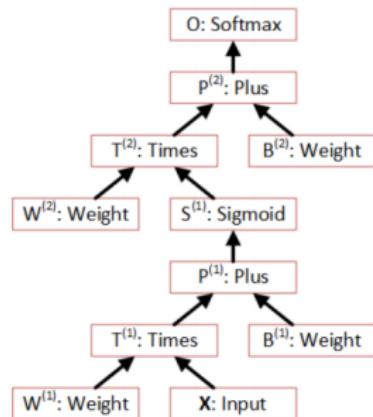
Frameworks

“Theano-like”

- Fine-grained
- Define **computational graphs** in a **symbolic** way
- Graph analysis and optimizations

Examples:

- Theano
- Computation Graph Toolkit (CGT)
- TensorFlow
- Computational Network Toolkit (CNTK)



(Kenneth Tran. “Evaluation of Deep Learning Toolkits”.

<https://github.com/zer0n/deepframeworks>)

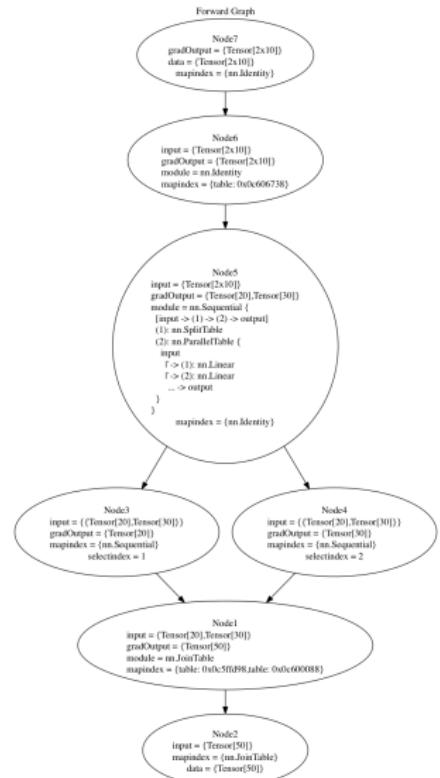
Frameworks

“Torch-like”

- Coarse-grained
- Build models by combining pre-specified modules
- Each module is manually implemented, hand-tuned

Examples:

- Torch7
- Caffe



Frameworks

Common in both:

- Define models using the framework's (constrained) symbolic language
- The framework handles backpropagation
→ you don't have to code derivatives
(unless adding new modules)
- Because derivatives are “automatic”, some call it “autodiff” or “automatic differentiation”

Frameworks

Common in both:

- Define models using the framework's (constrained) symbolic language
- The framework handles backpropagation
→ you don't have to code derivatives
(unless adding new modules)
- Because derivatives are “automatic”, some call it “autodiff” or “automatic differentiation”

This is NOT the traditional meaning of **automatic differentiation** (AD) (Griewank & Walther, 2008)

Frameworks

Common in both:

- Define models using the framework's (constrained) symbolic language
- The framework handles backpropagation
→ you don't have to code derivatives
(unless adding new modules)
- Because derivatives are “automatic”, some call it “autodiff” or “automatic differentiation”

This is NOT the traditional meaning of **automatic differentiation** (AD) (Griewank & Walther, 2008)

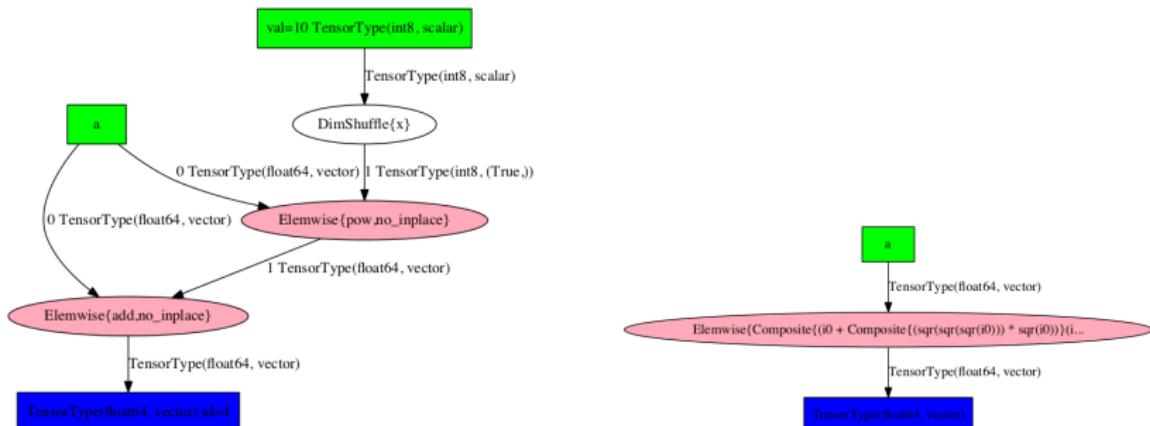
Because “automatic” is a generic (and bad) term,
algorithmic differentiation is a better name

“But, how is AD different from Theano?”

“But, how is AD different from Theano?”

In Theano

- express all math relations using symbolic placeholders
- use a **mini-language** with very **limited control flow** (e.g. scan)
- end up designing a symbolic graph for your algorithm
- Theano optimizes it



“But, how is AD different from Theano?”

Theano gives you automatic derivatives

- Transforms your graph into a derivative graph
- Applies optimizations
 - Identical subgraph elimination
 - Simplifications
 - Stability improvements
[\(http://deeplearning.net/software/theano/optimizations.html\)](http://deeplearning.net/software/theano/optimizations.html)
- Compiles to a highly optimized form

“But, how is AD different from Theano?”

You are limited to symbolic graph building, with the mini-language

"But, how is AD different from Theano?"

You are limited to symbolic graph building, with the mini-language

For example, instead of this in pure Python (for A^k):

```
result = 1
for i in xrange(k):
    result = result * A
```

"But, how is AD different from Theano?"

You are limited to symbolic graph building, with the mini-language

For example, instead of this in pure Python (for A^k):

```
result = 1
for i in xrange(k):
    result = result * A
```

You build this symbolic graph:

```
import theano
import theano.tensor as T

k = T.iscalar("k")
A = T.vector("A")

# Symbolic description of a Loop
result, updates = theano.scan(fn=lambda prior_result, A: prior_result * A,
                               outputs_info=T.ones_like(A),
                               non_sequences=A,
                               n_steps=k)

final_result = result[-1]

# compiled function that returns A**k
power = theano.function(inputs=[A,k], outputs=final_result, updates=updates)
```

"But, how is AD different from Theano?"

AD allows you to **just fully use your host language**
and gives you **exact and efficient derivatives**

"But, how is AD different from Theano?"

AD allows you to **just fully use your host language** and gives you **exact and efficient derivatives**

So, you just do this:

```
result = 1
for i in xrange(k):
    result = result * A
```

"But, how is AD different from Theano?"

AD allows you to **just fully use your host language** and gives you **exact and efficient derivatives**

So, you just do this:

```
result = 1
for i in xrange(k):
    result = result * A
```

For Python, autograd

<https://github.com/HIPS/autograd>

Harvard Intelligent Probabilistic Systems Group

(Dougal Maclaurin, David Duvenaud, Ryan P Adams. "Autograd: effortless gradients in Numpy." 2015)

Here is the difference

- AD does not use symbolic graphs
- Gives numeric code that **computes the function AND its derivatives** at a given point

```
f(a, b):           f'(a, a', b, b'): 
  c = a * b       (c, c') = (a*b, a'*b + a*b')
  d = sin c       (d, d') = (sin c, c' * cos c)
  return d         return (d, d')
```

- Derivatives propagated at the elementary operation level, as a side effect, at the same time when the function itself is computed
 - Prevents the “expression swell” of symbolic derivatives
- Full expressive capability of the host language
 - **Including conditionals, looping, branching**

Function evaluation traces

All **numeric evaluations** are sequences of elementary operations:
a “**trace**,” also called a “**Wengert list**” (Wengert, 1964)

```
f(a, b):
    c = a * b
    if c > 0
        d = log c
    else
        d = sin c
    return d
```

Function evaluation traces

All **numeric evaluations** are sequences of elementary operations:
a “**trace**,” also called a “**Wengert list**” (Wengert, 1964)

```
f(a, b):
    c = a * b
    if c > 0
        d = log c
    else
        d = sin c
    return d
```

```
f(2, 3)
```

Function evaluation traces

All **numeric evaluations** are sequences of elementary operations:
a “**trace**,” also called a “**Wengert list**” (Wengert, 1964)

```
f(a, b):          a = 2
    c = a * b
    if c > 0          b = 3
        d = log c
    else              c = a * b = 6
        d = sin c
    return d          d = log c = 1.791
```

```
f(2, 3)          return 1.791
```

(primal)

Function evaluation traces

All **numeric evaluations** are sequences of elementary operations:
a “**trace**,” also called a “**Wengert list**” (Wengert, 1964)

<code>f(a, b):</code>	<code>a = 2</code>	<code>a = 2</code>
<code>c = a * b</code>		
<code>if c > 0</code>	<code>b = 3</code>	<code>a' = 1</code>
<code>d = log c</code>		<code>b = 3</code>
<code>else</code>	<code>c = a * b = 6</code>	<code>b' = 0</code>
<code>d = sin c</code>		<code>c = a * b = 6</code>
<code>return d</code>	<code>d = log c = 1.791</code>	<code>c' = a' * b + a * b' = 3</code>
		<code>d = log c = 1.791</code>
		<code>d' = c' * (1 / c) = 0.5</code>
<code>f(2, 3)</code>	<code>return 1.791</code>	<code>return 1.791, 0.5</code>
	(primal)	(tangent)

Function evaluation traces

All **numeric evaluations** are sequences of elementary operations:
a “**trace**,” also called a “**Wengert list**” (Wengert, 1964)

<code>f(a, b):</code>	<code>a = 2</code>	<code>a = 2</code>
<code> c = a * b</code>		<code>a' = 1</code>
<code> if c > 0</code>	<code>b = 3</code>	<code>b = 3</code>
<code> d = log c</code>		<code>b' = 0</code>
<code> else</code>	<code>c = a * b = 6</code>	<code>c = a * b = 6</code>
<code> d = sin c</code>		<code>c' = a' * b + a * b' = 3</code>
<code> return d</code>	<code>d = log c = 1.791</code>	<code>d = log c = 1.791</code>
		<code>d' = c' * (1 / c) = 0.5</code>
<code>f(2, 3)</code>	<code>return 1.791</code>	<code>return 1.791, 0.5</code>
	(primal)	(tangent)

i.e., a Jacobian-vector product $\mathbf{J}_f(1, 0)|_{(2,3)} = \frac{\partial}{\partial a} f(a, b)|_{(2,3)} = 0.5$

This is called the **forward (tangent) mode** of AD

Function evaluation traces

```
f(a, b):
    c = a * b
    if c > 0
        d = log c
    else
        d = sin c
    return d
```

```
f(2, 3)
```

Function evaluation traces

```
f(a, b):           a = 2
                  c = a * b      b = 3
                  if c > 0      c = a * b = 6
                  d = log c      d = log c = 1.791
                  else
                  d = sin c
return d           return 1.791
                  (primal)
```

f(2, 3)

Function evaluation traces

```
f(a, b):  
    c = a * b  
    if c > 0  
        d = log c  
    else  
        d = sin c  
    return d
```

```
f(2, 3)
```

```
a = 2  
b = 3  
c = a * b = 6  
d = log c = 1.791  
return 1.791  
(primal)
```

```
a = 2  
b = 3  
c = a * b = 6  
d = log c = 1.791  
d' = 1  
c' = d' * (1 / c) = 0.166  
b' = c' * a = 0.333  
a' = c' * b = 0.5  
return 1.791, 0.5, 0.333
```

(adjoint)

Function evaluation traces

```
f(a, b):  
    c = a * b  
    if c > 0  
        d = log c  
    else  
        d = sin c  
    return d
```

```
f(2, 3)
```

```
a = 2  
b = 3  
c = a * b = 6  
d = log c = 1.791  
return 1.791  
(primal)
```

```
a = 2  
b = 3  
c = a * b = 6  
d = log c = 1.791  
d' = 1  
c' = d' * (1 / c) = 0.166  
b' = c' * a = 0.333  
a' = c' * b = 0.5  
return 1.791, 0.5, 0.333  
(adjoint)
```

i.e., a transposed Jacobian-vector product

$$\mathbf{J}_f^T(1)|_{(2,3)} = \nabla f|_{(2,3)} = (0.5, 0.333)$$

This is called the **reverse (adjoint) mode** of AD

Backpropagation is just a special case of the reverse mode:
code your neural network objective computation, apply reverse AD

Torch-autograd

There are signs that this type of **generalized AD** will become mainstream in machine learning

Torch-autograd

There are signs that this type of **generalized AD** will become mainstream in machine learning

A very recent development (November 2015)

Torch-autograd by Twitter Cortex

(inspired by Python autograd)

<https://blog.twitter.com/2015/autograd-for-torch>

“autograd has dramatically sped up our model building ... extremely easy to try and test out new ideas”

A cool functional DSL for Torch and Caffe

A side note about the **functional** interpretation deep learning:

dnngraph by Andrew Tulloch

<http://ajtulloch.github.io/dnngraph/>

Specify neural network layouts in Haskell,
it gives you Torch and Caffe scripts

What Can AD Research Contribute?

The ambition

- Deeply embedded AD
- Derivatives (forward and/or reverse) as part of the language infrastructure
- Rich API of differentiation operations as higher-order functions
- High-performance matrix operations for deep learning (GPU support, model and data parallelism)

The ambition

- Deeply embedded AD
- Derivatives (forward and/or reverse) as part of the language infrastructure
- Rich API of differentiation operations as higher-order functions
- High-performance matrix operations for deep learning (GPU support, model and data parallelism)

The embodiment of the “differentiable programming” paradigm

The ambition

- Deeply embedded AD
- Derivatives (forward and/or reverse) as part of the language infrastructure
- Rich API of differentiation operations as higher-order functions
- High-performance matrix operations for deep learning (GPU support, model and data parallelism)

The embodiment of the “differentiable programming” paradigm

I have been working on these issues with Barak Pearlmutter and created DiffSharp (later in the talk)

AD in a functional framework

AD has been around since the 1960s

(Wengert, 1964; Speelpenning, 1980; Griewank, 1989)

The foundations for AD in a functional framework

(Siskind and Pearlmutter, 2008; Pearlmutter and Siskind, 2008)

With research implementations

- R6RS-AD

<https://github.com/qobi/R6RS-AD>

- Stalingrad

<http://www.bcl.hamilton.ie/~qobi/stalingrad/>

- Alexey Radul's DVL

<https://github.com/axch/dysfunctional-language>

- Recently, my DiffSharp library

<http://diffsharp.github.io/DiffSharp/>

AD in a functional framework

"Generalized AD as a first-class function in an augmented λ-calculus" (Pearlmutter and Siskind, 2008)

Forward, reverse, and **any nested combination** thereof,
instantiated according to usage scenario

Nested lambda expressions with free-variable references

$$\min (\lambda x . (f x) + \min (\lambda y . g x y))$$

(min: gradient descent)

AD in a functional framework

“Generalized AD as a first-class function in an augmented λ-calculus” (Pearlmutter and Siskind, 2008)

Forward, reverse, and **any nested combination** thereof,
instantiated according to usage scenario

Nested lambda expressions with free-variable references

$$\min (\lambda x . (f x) + \min (\lambda y . g x y)) \\ (\text{min: gradient descent})$$

Must handle “perturbation confusion” (Manzyuk et al., 2012)

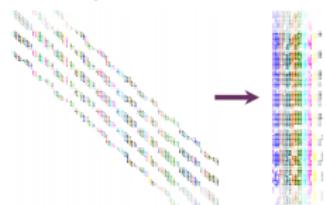
$$D (\lambda x . x \times (D (\lambda y . x + y) 1)) 1$$

$$\frac{d}{dx} \left(x \left(\frac{d}{dy} x + y \right) \Big|_{y=1} \right) \Bigg|_{x=1} \stackrel{?}{=} 1$$

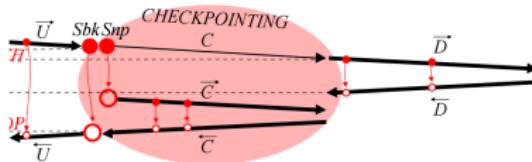
Tricks of the trade

Many methods from AD research

- Hessian-vector products (Pearlmutter, 1994)
- Tape reduction and elimination (Naumann, 2004)
- Context-aware source-to-source transformation (Utke, 2004)
- Utilizing sparsity by matrix coloring (Gebremedhin et al., 2013)



- Reverse AD checkpointing (Dauvergne & Hascoët, 2006)



My Ongoing Work

DiffSharp

<http://diffsharp.github.io/DiffSharp/>

- AD with linear algebra primitives
- arbitrary nesting of forward/reverse AD
- a comprehensive higher-order API
- gradients, Hessians, Jacobians, directional derivatives, matrix-free Hessian- and Jacobian-vector products



DiffSharp

<http://diffsharp.github.io/DiffSharp/>

- AD with linear algebra primitives
- arbitrary nesting of forward/reverse AD
- a comprehensive higher-order API
- gradients, Hessians, Jacobians, directional derivatives, matrix-free Hessian- and Jacobian-vector products

Implemented in F#

- the best tool for this job
- cross-platform (Linux, Mac OS, Windows)
- easy deployment with nuget
- the immense .NET user base of C# and F# users
- implicit quotations in F# 4.0 is a “killer feature” for deeply embedding transformation-based AD



DiffSharp

Higher-order differentiation API

Op.	Value	Type signature	AD	Num.	Sym.	
$f : \mathbb{R} \rightarrow \mathbb{R}$	f'	$(\mathbb{R} \rightarrow \mathbb{R}) \rightarrow \mathbb{R} \rightarrow \mathbb{R}$	X, F	A	X	
	f''	$(\mathbb{R} \rightarrow \mathbb{R}) \rightarrow \mathbb{R} \rightarrow (\mathbb{R} \times \mathbb{R})$	X, F	A	X	
	f'''	$(\mathbb{R} \rightarrow \mathbb{R}) \rightarrow \mathbb{R} \rightarrow \mathbb{R}$	X, F	A	X	
	(f, f'')	$(\mathbb{R} \rightarrow \mathbb{R}) \rightarrow \mathbb{R} \rightarrow (\mathbb{R} \times \mathbb{R})$	X, F	A	X	
	(f, f', f'')	$(\mathbb{R} \rightarrow \mathbb{R}) \rightarrow \mathbb{R} \rightarrow (\mathbb{R} \times \mathbb{R} \times \mathbb{R})$	X, F	A	X	
	$f^{(n)}$	$\mathbb{N} \rightarrow (\mathbb{R} \rightarrow \mathbb{R}) \rightarrow \mathbb{R} \rightarrow \mathbb{R}$	X, F		X	
	$f^{(n)}$	$\mathbb{N} \rightarrow (\mathbb{R} \rightarrow \mathbb{R}) \rightarrow \mathbb{R} \rightarrow (\mathbb{R} \times \mathbb{R})$	X, F		X	
$f : \mathbb{R}^n \rightarrow \mathbb{R}$	grad	∇f	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^n$	X, R	A	X
	grad'	$(f, \nabla f)$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R} \rightarrow (\mathbb{R} \times \mathbb{R}^n)$	X, R	A	X
	gradv	$\nabla f \cdot \mathbf{v}$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}$	X, F		A
	gradv'	$(f, \nabla f \cdot \mathbf{v})$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R} \times \mathbb{R})$	X, F		A
	hessian	\mathbf{H}_f	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^{n \times n}$	X, R-F	A	X
	$\text{hessian}'$	(f, \mathbf{H}_f)	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R} \times \mathbb{R}^{n \times n})$	X, R-F	A	X
	hessianv	$\mathbf{H}_f \mathbf{v}$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^n$	X, F-R	A	
	$\text{hessianv}'$	$(f, \mathbf{H}_f \mathbf{v})$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R} \times \mathbb{R}^n)$	X, F-R	A	
	gradhessian	$(\nabla f, \mathbf{H}_f)$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R}^n \times \mathbb{R}^{n \times n})$	X, R-F	A	X
	$\text{gradhessian}'$	$(f, \nabla f, \mathbf{H}_f)$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R} \times \mathbb{R}^n \times \mathbb{R}^{n \times n})$	X, R-F	A	X
	gradhessianv	$(\nabla f \cdot \mathbf{v}, \mathbf{H}_f \mathbf{v})$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R} \times \mathbb{R}^n)$	X, F-R	A	
	$\text{gradhessianv}'$	$(f, \nabla f \cdot \mathbf{v}, \mathbf{H}_f \mathbf{v})$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R} \times \mathbb{R} \times \mathbb{R}^n)$	X, F-R	A	
	laplacian	$\text{tr}(\mathbf{H}_f)$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}$	X, R-F	A	X
	$\text{laplacian}'$	$(f, \text{tr}(\mathbf{H}_f))$	$(\mathbb{R}^n \rightarrow \mathbb{R}) \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R} \times \mathbb{R})$	X, R-F	A	X
$f : \mathbb{R}^n \rightarrow \mathbb{R}^m$	jacobian	\mathbf{J}_f	$(\mathbb{R}^n \rightarrow \mathbb{R}^m) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^{m \times n}$	X, F/R	A	X
	$\text{jacobian}'$	(f, \mathbf{J}_f)	$(\mathbb{R}^n \rightarrow \mathbb{R}^m) \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R}^m \times \mathbb{R}^{m \times n})$	X, F/R	A	X
	jacobianv	$\mathbf{J}_f \mathbf{v}$	$(\mathbb{R}^n \rightarrow \mathbb{R}^m) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^m$	X, F		A
	$\text{jacobianv}'$	$(f, \mathbf{J}_f \mathbf{v})$	$(\mathbb{R}^n \rightarrow \mathbb{R}^m) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R}^m \times \mathbb{R}^m)$	X, F		A
	jacobianT	\mathbf{J}_f^T	$(\mathbb{R}^n \rightarrow \mathbb{R}^m) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^{n \times m}$	X, F/R	A	X
	$\text{jacobianT}'$	(f, \mathbf{J}_f^T)	$(\mathbb{R}^n \rightarrow \mathbb{R}^m) \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R}^m \times \mathbb{R}^{n \times m})$	X, F/R	A	X
	jacobianTv	$\mathbf{J}_f^T \mathbf{v}$	$(\mathbb{R}^n \rightarrow \mathbb{R}^m) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^m \rightarrow \mathbb{R}^n$	X, R		
	$\text{jacobianTv}'$	$(f, \mathbf{J}_f^T \mathbf{v})$	$(\mathbb{R}^n \rightarrow \mathbb{R}^m) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}^m \rightarrow (\mathbb{R}^m \times \mathbb{R}^n)$	X, R		
	$\text{jacobianTv}''$	$(f, \mathbf{J}_f^T(\cdot))$	$(\mathbb{R}^n \rightarrow \mathbb{R}^m) \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R}^m \times (\mathbb{R}^m \rightarrow \mathbb{R}^n))$	X, R		
	curl	$\nabla \times \mathbf{f}$	$(\mathbb{R}^3 \rightarrow \mathbb{R}^3) \rightarrow \mathbb{R}^3 \rightarrow \mathbb{R}^3$	X, F	A	X
	curl'	$(\mathbf{f}, \nabla \times \mathbf{f})$	$(\mathbb{R}^3 \rightarrow \mathbb{R}^3) \rightarrow \mathbb{R}^3 \rightarrow (\mathbb{R}^3 \times \mathbb{R}^3)$	X, F	A	X
	div	$\nabla \cdot \mathbf{f}$	$(\mathbb{R}^n \rightarrow \mathbb{R}^n) \rightarrow \mathbb{R}^n \rightarrow \mathbb{R}$	X, F	A	X
	div'	$(\mathbf{f}, \nabla \cdot \mathbf{f})$	$(\mathbb{R}^n \rightarrow \mathbb{R}^n) \rightarrow \mathbb{R}^n \rightarrow (\mathbb{R}^n \times \mathbb{R})$	X, F	A	X
curldiv	$(\nabla \times \mathbf{f}, \nabla \cdot \mathbf{f})$	$(\mathbb{R}^3 \rightarrow \mathbb{R}^3) \rightarrow \mathbb{R}^3 \rightarrow (\mathbb{R}^3 \times \mathbb{R})$	X, F	A	X	
	$\text{curldiv}'$	$(\mathbf{f}, \nabla \times \mathbf{f}, \nabla \cdot \mathbf{f})$	$(\mathbb{R}^3 \rightarrow \mathbb{R}^3) \rightarrow \mathbb{R}^3 \rightarrow (\mathbb{R}^3 \times \mathbb{R}^3 \times \mathbb{R})$	X, F	A	X

DiffSharp

Matrix operations

<http://diffsharp.github.io/DiffSharp/api-overview.html>

High-performance OpenBLAS backend by default, work on a CUDA-based GPU backend underway

Support for 64- and 32-bit floats (faster on many systems)

Benchmarking tool

<http://diffsharp.github.io/DiffSharp/benchmarks.html>

A growing collection of tutorials: gradient-based optimization algorithms, clustering, Hamiltonian Monte Carlo, neural networks, inverse kinematics

Hype

<http://hypelib.github.io/Hype/>

An experimental library for “compositional machine learning and **hyperparameter optimization**”, built on DiffSharp

A robust optimization core

- highly configurable functional modules
- SGD, conjugate gradient, Nesterov, AdaGrad, RMSProp, Newton's method
- Use nested AD for gradient-based hyperparameter optimization (Maclaurin et al., 2015)

Researching the differentiable functional programming paradigm for machine learning

Hype

Extracts from Hype neural network code,
use higher-order functions, don't think about gradients or
backpropagation

<https://github.com/hypelib/Hype/blob/master/src/Hype/Neural.fs>

```
1: // Use mixed mode nested AD
2: open DiffSharp.AD.Float32
3:
4: type FeedForward() =
5:     inherit Layer()
6:     // Feedforward layers executed as "fold", DM -> DM
7:     override n.Run(x:DM) = Array.fold Layer.run x layers
8:
9: type GRU(inputs:int, memcells:int) =
10:    inherit Layer()
11:    // RNN many-to-many execution as "map", DM -> DM
12:    override l.Run (x:DM) =
13:        x |> DM.mapCols
14:            (fun x ->
15:                let z = sigmoid(l.Wxz * x + l.Whz * l.h + l.bz)
16:                let r = sigmoid(l.Wxr * x + l.Whz * l.h + l.br)
17:                let h' = tanh(l.Wxh * x + l.Whh * (l.h .* r))
18:                l.h <- (1.f - z) .* h' + z .* l.h
19:                l.h)
```

Hype

Extracts from Hype optimization code

<https://github.com/hypelib/Hype/blob/master/src/Hype/Optimize.fs>

Optimization and training as higher-order functions

- works with any function that you want to describe your data
- can be composed, curried, nested

```
1: // Minimize function `f`  
2: static member Minimize (f:DV->D, w0:DV) =  
3:     Optimize.Minimize (f, w0, Params.Default)  
4:  
5: // Train model function `f`  
6: static member Train (f:DV->DV->D, w0:DV, d:Dataset) =  
7:     Optimize.Train ((fun w v -> toDV [f w v]), w0, d)
```

Hype

User doesn't need to think about derivatives

They are instantiated within the optimization code

```
1: type Method
2:     | CG -> // Conjugate gradient
3:         fun w f g p gradclip ->
4:             let v', g' = grad' f w // gradient
5:             let g' = gradclip g'
6:             let y = g' - g
7:             let b = (g' * y) / (p * y)
8:             let p' = -g' + b * p
9:             v', g', p'
10:    | NewtonCG -> // Newton conjugate gradient
11:        fun w f _ p gradclip ->
12:            let v', g' = grad' f w // gradient
13:            let g' = gradclip g'
14:            let hv = hessianv f w p // Hessian-vector product
15:            let b = (g' * hv) / (p * hv)
16:            let p' = -g' + b * p
17:            v', g', p'
18:    | Newton -> // Newton's method
19:        fun w f _ _ gradclip ->
20:            let v', g', h' = gradhessian' f w // gradient, Hessian
21:            let g' = gradclip g'
22:            let p' = -DM.solveSymmetric h' g'
23:            v', g', p'
```

Hype

But they can use derivatives within their models, if needed

- input sensitivities
- complex objective functions
- adaptive PID controllers
- integrating differential equations

```
1: // Leapfrog integrator, Hamiltonian
2: let leapFrog (u:DV->D) (k:DV->D) (d:D) steps (x0, p0) =
3:   let hd = d / 2.
4:   [1..steps]
5:   |> List.fold (fun (x, p) _ ->
6:     let p' = p - hd * grad u x
7:     let x' = x + d * grad k p'
8:     x', p' - hd * grad u x') (x0, p0)
```

Hype

But they can use derivatives within their models, if needed

- input sensitivities
- complex objective functions
- adaptive PID controllers
- integrating differential equations

```
1: // Leapfrog integrator, Hamiltonian
2: let leapFrog (u:DV->D) (k:DV->D) (d:D) steps (x0, p0) =
3:   let hd = d / 2.
4:   [1..steps]
5:   |> List.fold (fun (x, p) _ ->
6:     let p' = p - hd * grad u x
7:     let x' = x + d * grad k p'
8:     x', p' - hd * grad u x') (x0, p0)
```

Thanks to nested generalized AD

- you can optimize components that are internally using differentiation
- resulting higher-order derivatives propagate via forward/reverse AD as needed

Hype

We also provide a Torch-like API for neural networks

```
1: let n = FeedForward()  
2: n.Add(Linear(dim, 100))  
3: n.Add(LSTM(100, 400))  
4: n.Add(LSTM(400, 100))  
5: n.Add(Linear(100, dim))  
6: n.Add(reLU)
```

Hype

We also provide a Torch-like API for neural networks

```
1: let n = FeedForward()  
2: n.Add(Linear(dim, 100))  
3: n.Add(LSTM(100, 400))  
4: n.Add(LSTM(400, 100))  
5: n.Add(Linear(100, dim))  
6: n.Add(ReLU)
```

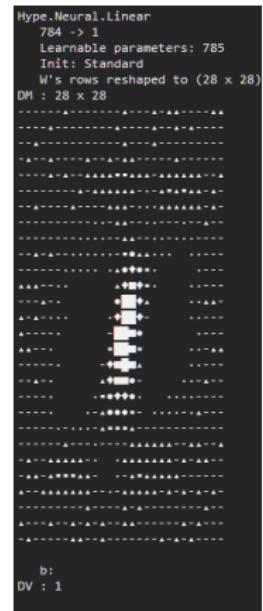
A cool thing: thanks to AD, we can freely code
any F# function as a layer, it just works

```
1: n.Add(fun m -> m |> DM.mapCols softmax) // A "map" of softmax  
2:  
3: let dropout (x:DM) = // Implement a new layer (dropout)  
4:     x .* (Rnd.UniformDM(x.Cols, x.Rows) |> DM.Round) * 2.f  
5:  
6: n.Add(dropout) // Add any function as a layer
```

Hype

<http://hypelib.github.io/Hype/feedforwardnets.html>

We also have some nice additions for F# interactive



Roadmap

- Transformation-based, context-aware AD
F# quotations (Syme, 2006) give us a direct path for deeply embedding AD
- Currently experimenting with GPU backends
(CUDA, ArrayFire, Magma)
- Generalizing to tensors
(for elegant implementations of, e.g., ConvNets)

Roadmap

I would like to see this work integrated with tools in other languages (C++, Python) and frameworks (Torch, CNTK)

Conclusion

Conclusion

An exciting research area at the intersection of

- programming languages
- functional programming
- machine learning

Beyond deep learning

Applications in probabilistic programming

(Wingate, Goodman, Stuhlmüller, Siskind. "Nonstandard interpretations of probabilistic programs for efficient inference." 2011)

- Hamiltonian Monte Carlo

[http://diffsharp.github.io/DiffSharp/
examples-hamiltonianmontecarlo.html](http://diffsharp.github.io/DiffSharp/examples-hamiltonianmontecarlo.html)

- No-U-Turn sampler

- Gradient-based maximum a posteriori estimates

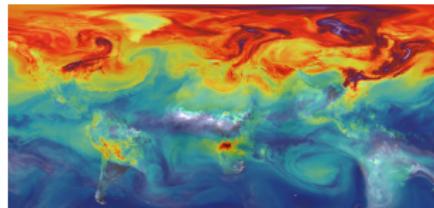
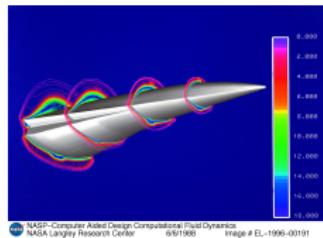
For example, Stan is built on AD

<http://mc-stan.org/>
(Carpenter et al., 2015)

Other areas

Any work in AD remains applicable to the
traditional application domains of AD in industry and academia
(Corliss et al., 2002)

- Computational fluid dynamics
- Atmospheric chemistry
- Engineering design optimization
- Computational finance



Thank You!

References

- Baydin AG, Pearlmutter BA, Radul AA, Siskind JM (Submitted) Automatic differentiation in machine learning: a survey [arXiv:1502.05767]
- Baydin AG, Pearlmutter BA, Siskind JM (Submitted) DiffSharp: automatic differentiation library [arXiv:1511.07727]
- Bengio Y (2013) Deep learning of representations: looking forward. Statistical Language and Speech Processing. LNCS 7978:1–37 [arXiv:1404.7456]
- Graves A, Wayne G, Danihelka I (2014) Neural Turing machines. [arXiv:1410.5401]
- Grefenstette E, Hermann KM, Suleyman M, Blunsom, P (2015) Learning to transduce with unbounded memory. [arXiv:1506.02516]
- Griewank A, Walther A (2008) Evaluating Derivatives: Principles and Techniques of Algorithmic Differentiation. Society for Industrial and Applied Mathematics, Philadelphia [DOI 10.1137/1.9780898717761]
- He K, Zhang X, Ren S, Sun J (2015) Deep residual learning for image recognition. [arXiv:1512.03385]
- Joulin A, Mikolov T (2015) Inferring algorithmic patterns with stack-augmented recurrent nets. [arXiv:1503.01007]
- Maclaurin D, David D, Adams RP (2015) Gradient-based Hyperparameter Optimization through Reversible Learning [arXiv:1502.03492]
- Manzyuk O, Pearlmutter BA, Radul AA, Rush DR, Siskind JM (2012) Confusion of tagged perturbations in forward automatic differentiation of higher-order functions [arXiv:1211.4892]
- Pearlmutter BA, Siskind JM (2008) Reverse-mode AD in a functional framework: Lambda the ultimate backpropagator. ACM TOPLAS 30(2):7 [DOI 10.1145/1330017.1330018]
- Siskind JM, Pearlmutter BA (2008) Nesting forward-mode AD in a functional framework. Higher Order and Symbolic Computation 21(4):361–76 [DOI 10.1007/s10990-008-9037-1]
- Sukhbaatar S, Szlam A, Weston J, Fergus R (2015) Weakly supervised memory networks. [arXiv:1503.08895]
- Syme D (2006) Leveraging .NET meta-programming components from F#: integrated queries and interoperable heterogeneous execution. 2006 Workshop on ML. ACM.
- Vinyals O, Toshev A, Bengio S, Erhan D (2014) Show and tell: a neural image caption generator. [arXiv:1411.4555]
- Wengert R (1964) A simple automatic derivative evaluation program. Communications of the ACM 7:463–4
- Zaremba W, Sutskever I (2015) Reinforcement learning neural Turing machines. [arXiv:1505.00521]